

Contact

www.linkedin.com/in/keenan-venuti-in (LinkedIn)
keenanvenuti.com/home
(Personal)

Top Skills

Machine Learning
Python
Google Cloud Platform (GCP)

Certifications

Azure Fundamentals
Neo4j Certified Professional
Tensorflow Developer
Machine Learning

Keenan Venuti

Machine Learning Engineer
San Francisco, California, United States

Summary

Passionate about Automation, AI, Robotics, Cloud Technology, and Machine Learning from both a technical and philosophical perspective. To learn more about me, my projects, and my papers, visit keenanvenuti.com.

In my spare time, I'm an e-learning consumer, advocative philosopher, preachy evolutionist, pickup basketball MVP, and rap enthusiast.

I'm a certified TensorFlow Developer, Google Cloud Data Engineer, Azure, AWS, Stanford ML, and Neo4j Professional.

Experience

BARK

Lead Machine Learning Engineer
February 2021 - Present (2 years 3 months)
San Francisco Bay Area

(promoted from Sr. MLE)

- Owner of Bark's customer facing Add to Box recommender system and several estimated delivery date models.
- Currently developing an internal ML platform equipped with a feature store, automated training pipelines with hyperparameter tuning, and a microservices based inference layer with scalable endpoints and prediction logging. Implemented with GCP's core data engineering and machine learning products.
- Engineering lead in migrating Bark's production ML models (including those not listed above) from disparate deployments to our standardized internal ML platform.

Accenture

Data Science
September 2019 - February 2021 (1 year 6 months)
San Francisco Bay Area

-- Machine Learning Engineer --

- Lead ML Engineer for an NLU Question-Answering search capability for +30m life sciences abstract database.
- Assisted development of a massive (multi-million edges/nodes) life sciences research Knowledge Graph hosted on Google Cloud Platform.

-- Data Science Practitioner --

- Data Scientist for a large Healthcare provider. Developed ETL processes for a 50k member dataset across 5 databases and designed a distance algorithm to calculate member similarity across binary, categorical, and quantitative dimensions.
 - Engineered unsupervised ML methods (clustering, PFA) to programmatically identify potential member groupings and important features. Presented my work to the client's Digital SVP and several VPs.
 - Developed an automated document extraction and processing pipeline on Google Cloud Platform (GCP) using Natural Language Processing and GCP's ML resources for a large Financial Services client
 - Predicted gaps in care for patients with diabetes and high blood pressure using R for a large Health Insurance client
- Leadership Positions --
- U.S. Data Science Lead for INTIENT Research, a multi-million dollar life sciences platform.
 - Lead team of 3 building tools for Graph Database ingestion, harmonization, visualization, analytics, and machine learning.

Adobe

Data Analytics

January 2018 - September 2019 (1 year 9 months)

San Francisco Bay Area

- Developed automated ETL workflows in Alteryx and Snowflake to clean, aggregate, dissect, and publish product usage and sales data to Tableau. These data are consistently used in Adobe integration planning
- Data analyst for specialized data pulls across the company. Insights used in executive presentations, product revenue modeling, conferences, and other team projects
- Built data science library in Python for customer clustering, pricing, and sales hypothesis testing using Snowflake SQL Pipeline
- Built and maintained Tableau dashboards on product performance, win rate, competitive insights, and other metrics. Supports Product Management, Customer, and Marketing teams

- Modeled multiple regressions in SkLearn to predict server usage from customer purchase patterns
 - Programmed an online Marketo Pricing calculator in HTML and Javascript as internal Google Web App
- *Adobe acquired Marketo Oct. 2018*

Harvard Business School
 Researcher for HBS - Dr. Norton Lab
 August 2016 - December 2016 (5 months)
 Cambridge, MA

- Performed data collection and classification
- Wrote scripts in Python to parse unformatted text data into CSVs and encoded based on sentiment and theme
- Researched and authored a study proposal applying regression analyses to digital advertising headlines bucketed by demographic segments

Oracle
 Content Metrics (Intern)
 May 2016 - August 2016 (4 months)
 Redwood Shores

- Created a content metrics model using raw performance input data to output trend analyses and KPIs for marketing team to measure overall lift and find top performing posts, authors, and blogs
- Aligned marketing department with Oracle Labs and Oracle's Demand Generation team to create content and reformat their content repository database

Sharethrough
 Research (Intern)
 May 2015 - August 2015 (4 months)
 San Francisco Bay Area

- Envisioned and planned a headline-text analysis project. QUARTZ reported on the project: "These are the words people can't resist in a headline"
- Processed textual data using linguistic analysis program (LIWC) with research team

Sharethrough
 Engagement (Intern)
 June 2014 - December 2014 (7 months)

San Francisco Bay Area

- Parsed and cleaned raw advertising performance data for quarterly reporting
- Created sales enablement decks, leading to a record breaking quarter in sales, based on year to date data pulled from Sharethorough's exchange
- Managed and created several advertising campaigns on Twitter, Facebook, LinkedIn, and AdRoll

Education

Harvard University

Bachelor's degree, Psychology & Computer Science

Northwestern University

Master of Science - MS, Data Science, Artificial Intelligence Specialization

Gunn High School